

Methods of Knowledge Based Software Development /

Teadmispõhise tarkvaraarenduse meetodid

Tanel Tammet, Juhan Ernits

Department of Computer Science

Tallinn University of Technology

tanel.tammet@ttu.ee, juhan.ernits@ttu.ee

2016

Organisation of the course

- Lectures: Mondays 14:00 CYB-VEENUS
- Lab sessions: Fridays 14:00 ICT-401
- Course web page: <http://courses.cs.ttu.ee/pages/ITI8600>
- 4 hand in problems each worth of 10% of exam mark. Involves coding!
 - Some problems will be split in 2 subtasks.
 - Missing deadlines will lose you 1 point per day! Try to be on time! (If you have welfare reasons for delays, please let me know as they emerge, you can get deadline extensions, but upon prior agreement).
- Written exam. Worth 60% of the mark!
 - NB! Make sure to do the home assignments, otherwise your maximum will be “0”!
 - You will have to have at least 5 points for each home assignment and have submitted at least 3 homeworks to get to the exam.
- Goes without saying: try to solve the hand in problems yourself. You will have to explain what the code does upon request. If multiple solutions look too much alike, all involved parties will get 0% for that problem.

Hand in problems

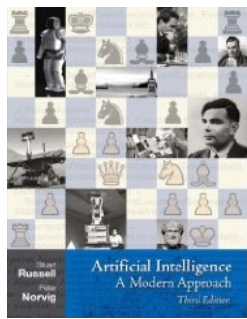
- Involve programming.
 - Main language used in the course is Python, but C++ and F# are encouraged as well. Yes, Java is OK too!
- It is important that you are able to encode the problem in such a way that you can solve it by writing a program!
- Pair programming will be part of the learning process.
- Submission to university GIT repositories (no e-mails)

Questions and answers sessions

- There will be a limited number of Q&A sessions offered so that you will be able to come and ask questions that have arisen during reading the links /the book or solving homework problems.
- If you feel that you cannot even formulate a question, come to the sessions!

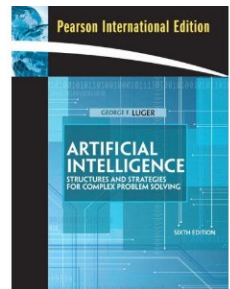
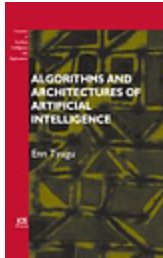
Exam

- 60% of the final mark
- You are asked to solve problems from various topics covered in the course on paper.
 - It can be quite hard, but the homeworks are meant to help you



Literature

- Web resources.
- Main book (for search and learning modules):
 - Russell and Norvig. Artificial Intelligence: A Modern Approach. 3rd edition 2010. http://tallinn.ester.ee/record=b2881231~S1*est
http://tallinn.ester.ee/record=b3000919~S1*est
- Additional reading:
 - Enn Tõugu. Algorithms and architectures of artificial intelligence. IOS Press, Amsterdam 2007.
http://tallinn.ester.ee/record=b2291064~S1*est
 - George F Luger. Artificial Intelligence: Structures and Strategies for complex problem solving. 6th edition, 2009.
http://tallinn.ester.ee/record=b2881423~S1*est
- This and that from other sources too (check the web page).



Perspective

- The course covers many topics – could be split into several courses, but:
 - Machine learning is addressed in depth in another course taught during the Spring semester.
 - Logic can be pursued in a separate logics course
 - Data mining is a separate course
- Too few copies of the book at the library:
 - You can also lease the AIMA book electronically at <http://www.mypearsonstore.com/bookstore/artificial-intelligence-a-modern-approach-coursesmart-9780136067337>.

Methods of Knowledge Based Software Development

Main topics of AI

- Natural language processing
- Knowledge representation
- Automated reasoning
- Machine learning
- Computer vision
- Robotics

Topics covered in other courses

- Robotics: IBX0020 – Robotics (Maarja Kruusmaa)
- Machine learning: ITI8565 Machine learning (Sven Nõmm)
- Natural Language Processing (thesis projects available from Dept. of Computer Science and Institute of Cybernetics (Tanel Alumäe, Einar Meister))
- Logic for Computer Science (Tarmo Uustalu)
- Data Mining (Innar Liiv and Sven Nõmm)

State of the art

- Deep Blue defeated the reigning world chess champion Garry Kasparov in 1997
- Proved a mathematical conjecture (Robbins conjecture) unsolved for decades
- No hands across America (driving autonomously 98% of the time from Pittsburgh to San Diego)
- 100000s of miles of autonomous driving (Google car)
- IBM DeepBlue played Chess at grandmaster level
- NASA's on-board autonomous planning program controlled the scheduling of operations for a spacecraft
- Proverb solves crossword puzzles better than most humans
- IBM Watson won in Jeopardy! in 2011.
- Medical diagnosis (e.g. diagnosis of diabetes)
- Robotics: Mars rovers, underwater robots, factories
- Learning to play ATARI games with deep reinforcement learning
- Watson, Google Now, Cordana, Siri, ...

FEATURES

Uber's First Self-Driving Fleet Arrives in Pittsburgh This Month

The autonomous cars, launching this summer, are custom Volvo XC90s, supervised by humans in the driver's seat.

August 18, 2016
by Max Chafin
from BloombergBusinessweek

Subscribe | Reprints



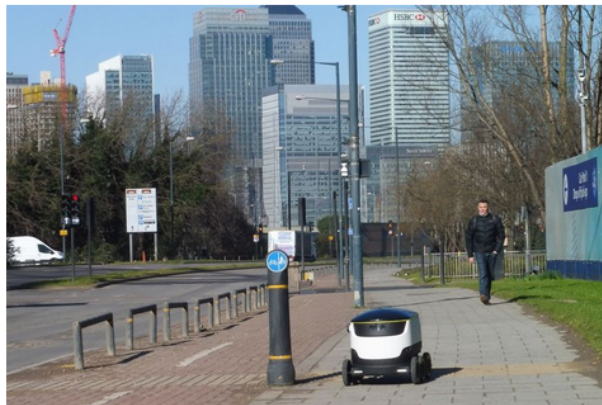
Robots

Starship partners with Just Eat to bring its delivery robots to the UK

The firm has also signed deals with Hermes, Metro Group, and startup Pronto.co.uk

By **VICTORIA WOOLLASTON**

05 Jul 2016



A peek into ethical questions

- AI has the same ethical problems as other, conventional artifacts.
 - Understanding what a realistic experience of AI is, is likely to help us better judge what it means to be human
- Developers of AI have an obligation NOT to exploit people's ignorance and make them think AI is human
- Robots are not really your friends
 - We build them with certain goals. Both buyers and builders should pick goals sensibly

Tesla

Tesla driver killed while using autopilot was watching Harry Potter, witness says

Driver in first known fatal self-driving car crash was also driving so fast that 'he went so fast through my trailer I didn't see him', the truck driver involved said



Investigation launched after Tesla driver killed while using autopilot

Solving problems by searching

Outline

- Problem-solving agents
- Problem types
- Problem formulation
- Example problems
- Basic search algorithms

Problem-solving agents

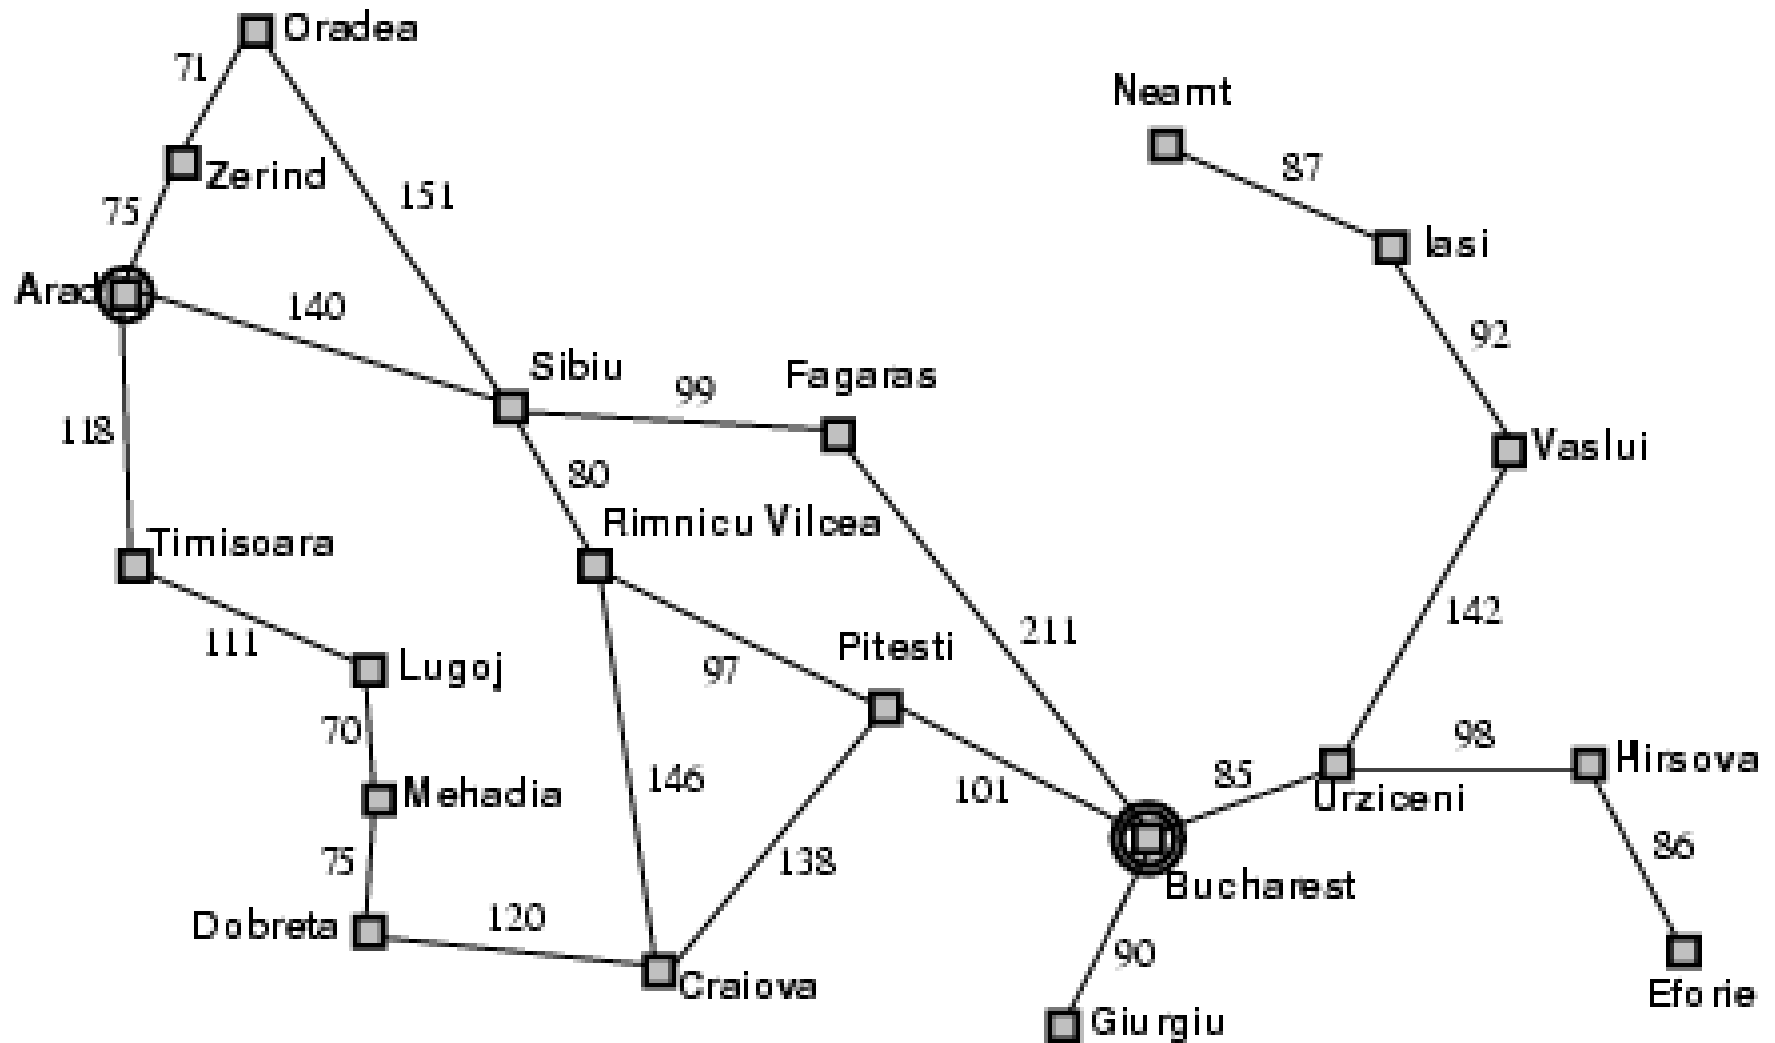
```
function SIMPLE-PROBLEM-SOLVING-AGENT(percept) returns an action
  static: seq, an action sequence, initially empty
           state, some description of the current world state
           goal, a goal, initially null
           problem, a problem formulation

  state  $\leftarrow$  UPDATE-STATE(state, percept)
  if seq is empty then do
    goal  $\leftarrow$  FORMULATE-GOAL(state)
    problem  $\leftarrow$  FORMULATE-PROBLEM(state, goal)
    seq  $\leftarrow$  SEARCH(problem)
  action  $\leftarrow$  FIRST(seq)
  seq  $\leftarrow$  REST(seq)
  return action
```

Example: Romania

- On holiday in Romania; currently in Arad.
- Flight leaves tomorrow from Bucharest
- **Formulate goal:**
 - be in Bucharest
- **Formulate problem:**
 - **states:** various cities
 - **actions:** drive between cities
- **Find solution:**
 - sequence of cities, e.g., Arad, Sibiu, Fagaras, Bucharest

Example: Romania

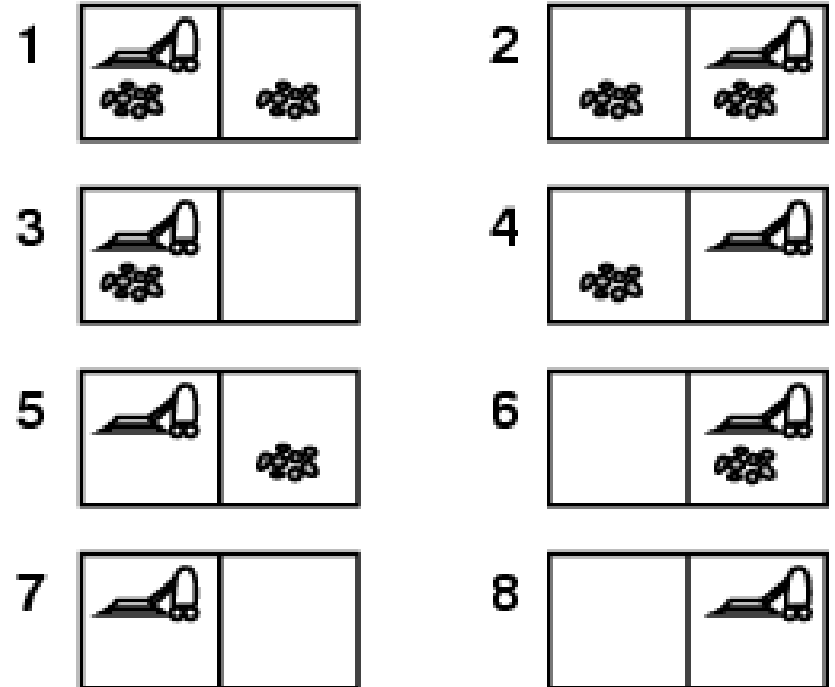


Problem types

- Deterministic, fully observable → single-state problem
 - Agent knows exactly which state it will be in; solution is a sequence
- Non-observable → sensorless problem (conformant problem)
 - Agent may have no idea where it is; solution is a sequence
- Nondeterministic and/or partially observable → contingency problem
 - percepts provide new information about current state
 - often interleave search and execution
- Unknown state space → exploration problem

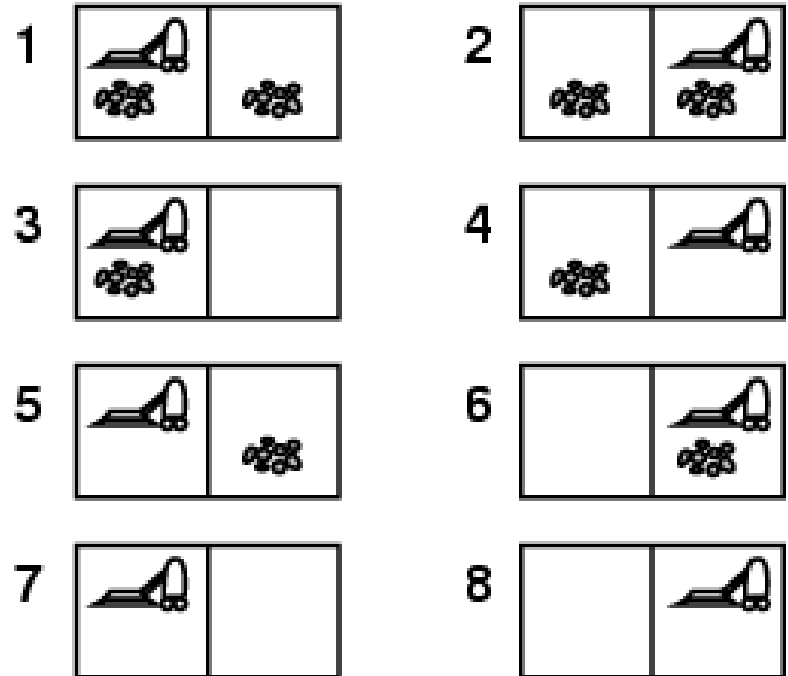
Example: vacuum world

- Single-state, start in #5.
Solution?



Example: vacuum world

- **Single-state**, start in #5.
Solution? [*Right, Suck*]
- **Sensorless**, start in {1,2,3,4,5,6,7,8} e.g.,
Right goes to {2,4,6,8}
Solution?



Example: vacuum world

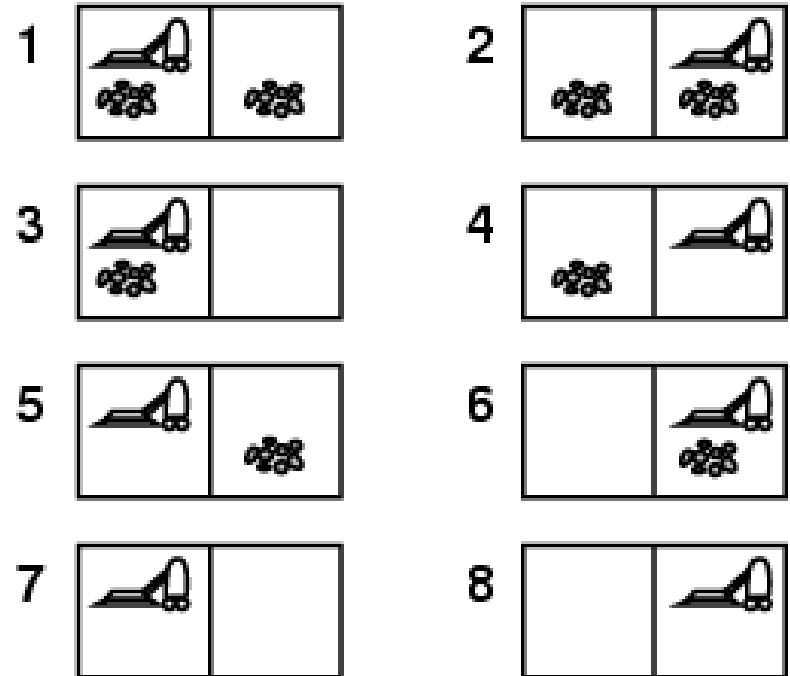
- **Sensorless**, start in {1,2,3,4,5,6,7,8} e.g.,
Right goes to {2,4,6,8}

Solution?

[Right, Suck, Left, Suck]

- **Contingency**
 - Nondeterministic: *Suck* may dirty a clean carpet
 - Partially observable: location, dirt
 - Percept: *[L, Clean]*, i.e., start in #5 or #7

Solution?



Example: vacuum world

- **Sensorless**, start in {1,2,3,4,5,6,7,8} e.g.,
Right goes to {2,4,6,8}

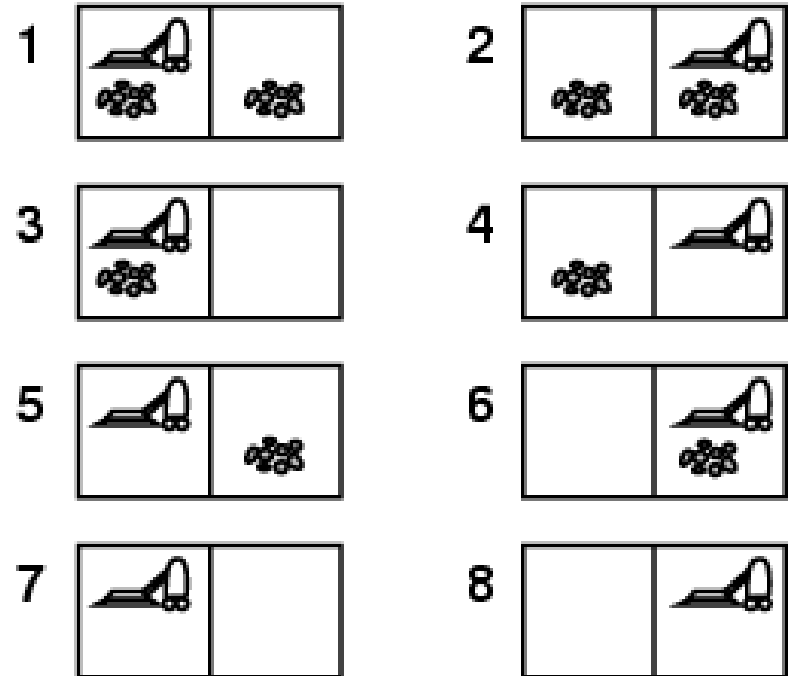
Solution?

[Right, Suck, Left, Suck]

- **Contingency**

- Nondeterministic: *Suck* may dirty a clean carpet
- Partially observable: location, dirt at current location.
- Percept: *[L, Clean]*, i.e., start in #5 or #7

Solution? *[Right, **if** dirt **then** Suck]*



Single-state problem formulation

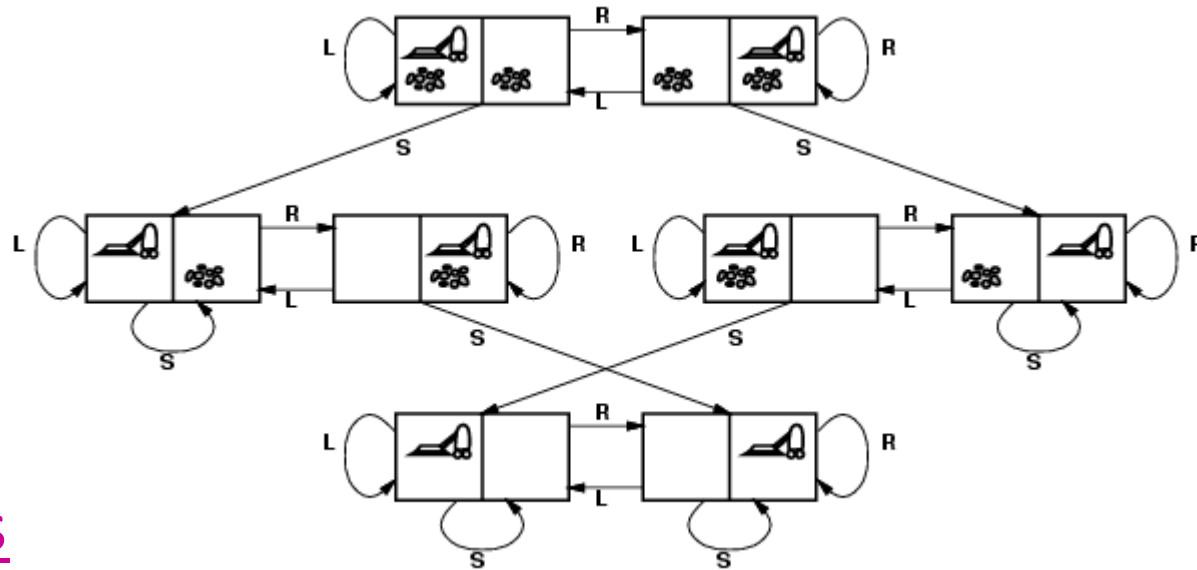
A **problem** is defined by four items:

1. **initial state** e.g., "at Arad"
 2. **actions** or **successor function** $S(x)$ = set of action–state pairs
 - e.g., $S(\text{Arad}) = \{ \langle \text{Arad} \rightarrow \text{Zerind}, \text{Zerind} \rangle, \dots \}$
 3. **goal test**, can be
 - **explicit**, e.g., $x = \text{"at Bucharest"}$
 - **implicit**, e.g., $\text{Checkmate}(x)$
 4. **path cost** (additive)
 - e.g., sum of distances, number of actions executed, etc.
 - $c(x, a, y)$ is the **step cost**, assumed to be ≥ 0
- A **solution** is a sequence of actions leading from the initial state to a goal state

Selecting a state space

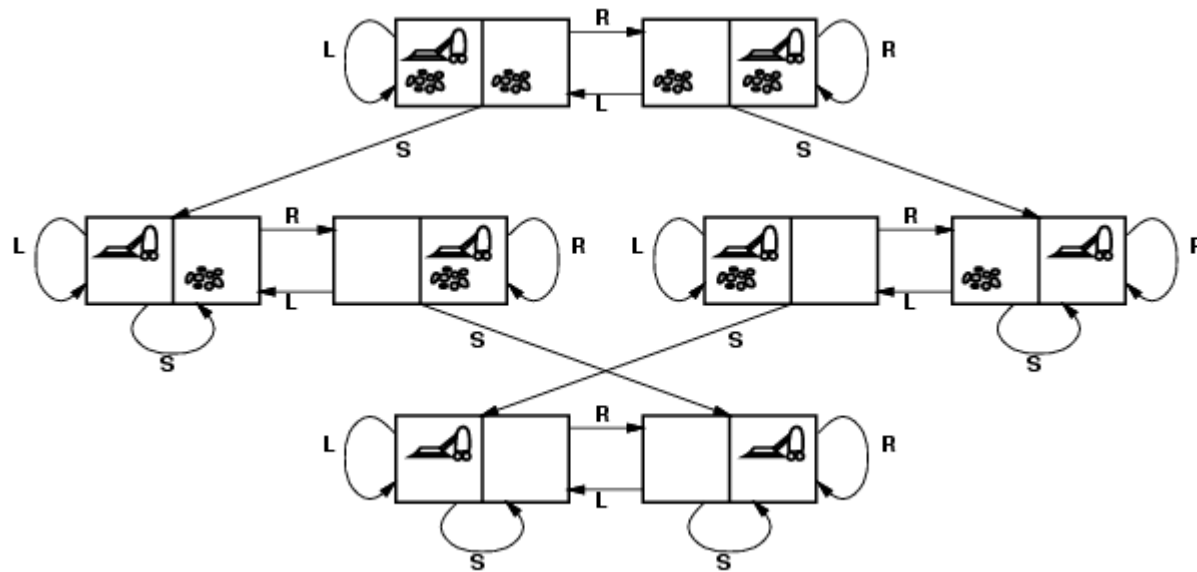
- Real world is absurdly complex
 - state space must be **abstracted** for problem solving
- (Abstract) state = set of real states
- (Abstract) action = complex combination of real actions
 - e.g., "Arad → Zerind" represents a complex set of possible routes, detours, rest stops, etc.
- For guaranteed realizability, **any** real state "in Arad" must get to **some** real state "in Zerind"
- (Abstract) solution =
 - set of real paths that are solutions in the real world
- Each abstract action should be "easier" than the original problem

Vacuum world state space graph



- states
- actions?
- goal test?
- path cost?
-

Vacuum world state space graph



- states? integer dirt and robot location
- actions? *Left, Right, Suck*
- goal test? no dirt at all locations
- path cost? 1 per action

Example: The 8-puzzle

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

- states?
- actions?
- goal test?
- path cost?

Example: The 8-puzzle

7	2	4
5		6
8	3	1

Start State

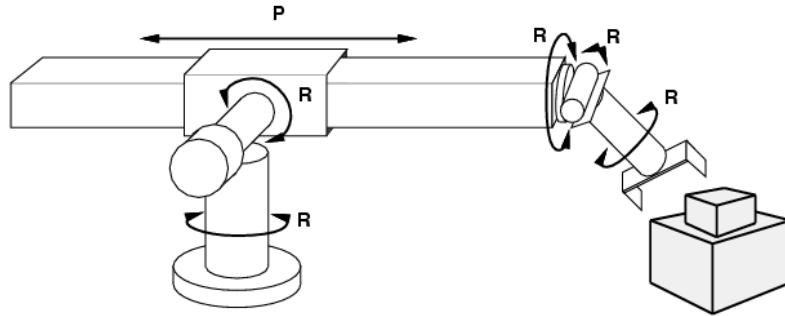
	1	2
3	4	5
6	7	8

Goal State

- states? locations of tiles
- actions? move blank left, right, up, down
- goal test? = goal state (given)
- path cost? 1 per move

[Note: optimal solution of n -Puzzle family is NP-hard]

Example: robotic assembly



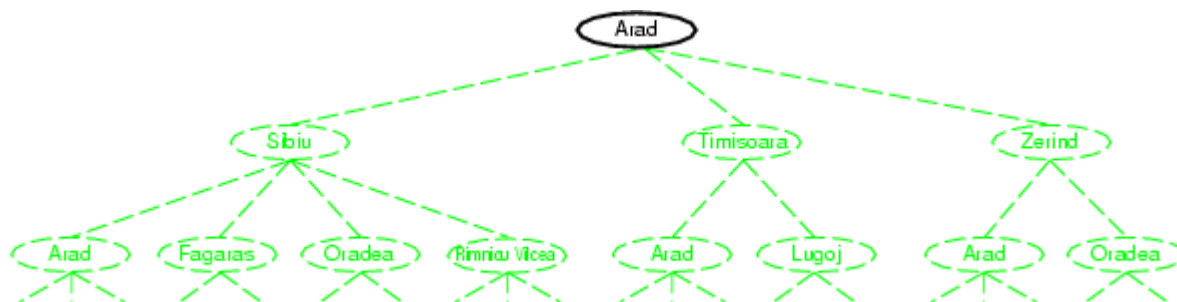
- states?: real-valued coordinates of robot joint angles
parts of the object to be assembled
- actions?: continuous motions of robot joints
- goal test?: complete assembly
- path cost?: time to execute

Tree search algorithms

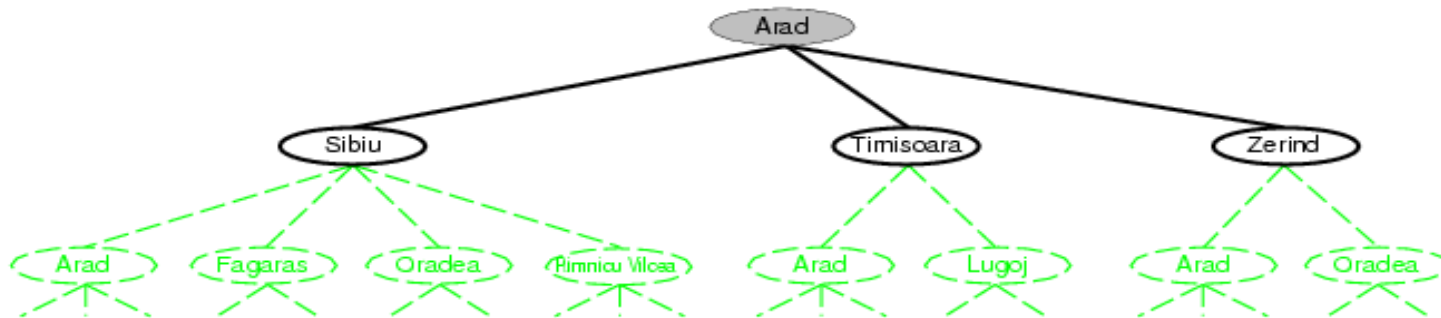
- Basic idea:
 - offline, simulated exploration of state space by generating successors of already-explored states (a.k.a. **expanding** states)

```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion then return failure
    choose a leaf node for expansion according to strategy
    if the node contains a goal state then return the corresponding solution
    else expand the node and add the resulting nodes to the search tree
```

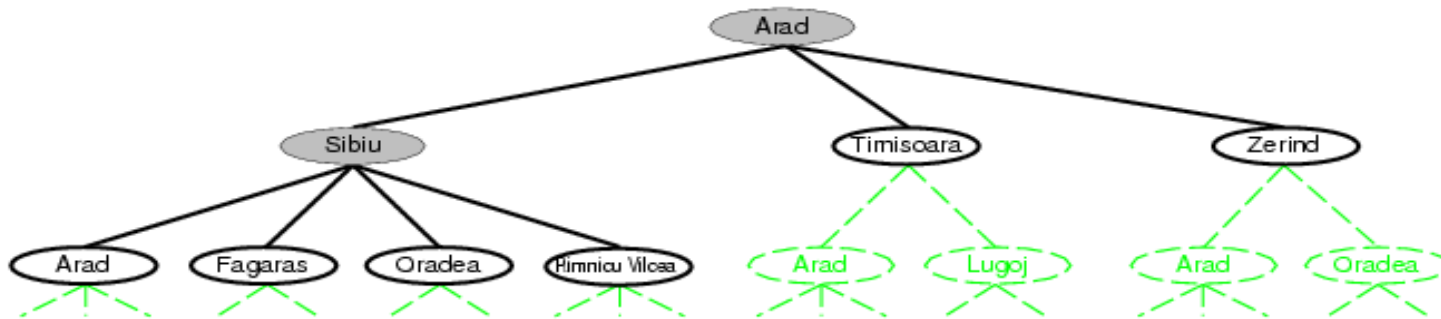
Tree search example



Tree search example



Tree search example



Implementation: general tree search

```
function TREE-SEARCH(problem frontier) returns a solution, or failure  
  frontier  $\leftarrow$  INSERT(MAKE-NODE(INITIAL-STATE[problem]), frontier)  
  loop do  
    if frontier empty then return failure  
    node  $\leftarrow$  REMOVE-FRONT(frontier)  
    if GOAL-TEST[problem](STATE[node]) then return SOLUTION(node)  
    frontier  $\leftarrow$  INSERTALL(EXPAND(node, problem), frontier)
```

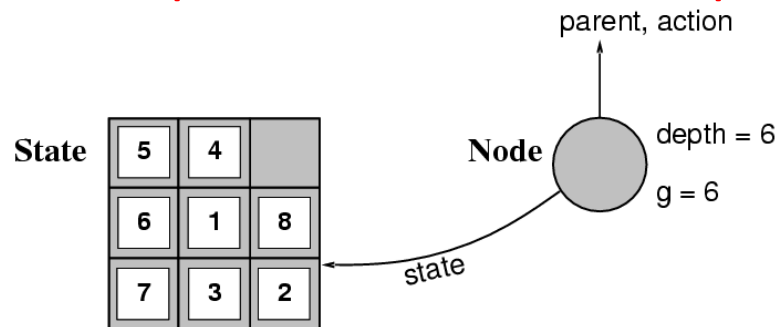
```
function EXPAND(node, problem) returns a set of nodes  
  successors  $\leftarrow$  the empty set  
  for each action, result in SUCCESSOR-FN[problem](STATE[node]) do  
    s  $\leftarrow$  a new NODE  
    PARENT-NODE[s]  $\leftarrow$  node; ACTION[s]  $\leftarrow$  action; STATE[s]  $\leftarrow$  result  
    PATH-COST[s]  $\leftarrow$  PATH-COST[node] + STEP-COST(node, action, s)  
    DEPTH[s]  $\leftarrow$  DEPTH[node] + 1  
    add s to successors  
  return successors
```

Tree search in Python

```
def tree_search(problem, frontier):  
    """Search through the successors of a problem to find a goal.  
    The argument frontier should be an empty queue.  
    Don't worry about repeated paths to a state. [Fig. 3.7]"""  
    frontier.append(Node(problem.initial))  
    while frontier:  
        node = frontier.pop()  
        if problem.goal_test(node.state):  
            return node  
        frontier.extend(node.expand(problem))  
    return None
```

Implementation: states vs. nodes

- A **state** is a (representation of) a physical configuration
- A **node** is a data structure constituting part of a search tree includes **state**, **parent node**, **action**, **path cost $g(x)$** , **depth**



- The `Expand` function creates new nodes, filling in the various fields and using the `SuccessorFn` of the problem to create the corresponding states.

Graph search

```
function GRAPH-SEARCH(problem, fringe) returns a solution, or failure
  closed ← an empty set
  fringe ← INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node ← REMOVE-FRONT(fringe)
    if GOAL-TEST[problem](STATE[node]) then return SOLUTION(node)
    if STATE[node] is not in closed then
      add STATE[node] to closed
      fringe ← INSERTALL(EXPAND(node, problem), fringe)
```

Frontier is sometimes called *fringe*.

Graph search in Python

```
def graph_search(problem, frontier):  
    """Search through the successors of a problem to find a goal.  
    The argument frontier should be an empty queue.  
    If two paths reach a state, only use the first one. [Fig. 3.7]"""  
    frontier.append(Node(problem.initial))  
    explored = set()  
    while frontier:  
        node = frontier.pop()  
        if problem.goal_test(node.state):  
            return node  
        explored.add(node.state)  
        frontier.extend(child for child in node.expand(problem)  
                        if child.state not in explored  
                        and child not in frontier)  
    return None
```

Acknowledgements

- This set of slides contains several prepared by Hwee Tou Ng and Stuart Russell, available from [the AIMA pages](#).